



Query Store and Azure SQL Copilot, who is the fairest in the land?

Jose Manuel Jurado

Escalation Engineer Microsoft

- **Services Delivery Excellence Team for PaaS databases.**
- 14 Years @ Microsoft
- **Supported SQL Server Core, Analysis Services and SQL Server On Azure VMs or other RDBMS.**
- Working with DevOps (ARM), startups companies, Developers, DBA, DBM, CIO, CFO, etc.

-
- **Speaker @SQL Saturday, TechReady, SQL Nexus, SQLBits, Microsoft Summit, Azure Global BootCamp, SQLKofenrenz, SQLDay.**
 - Worked previously as DBA, Developer, IT Manager and other jobs for more than 25 years which I don't even remember anymore.
Microsoft Certified Trainer and other certifications.



/josemanueljuradodiaz



@jmJuradoDiaz



José Manuel Jurado Díaz



Jmjurado@Microsoft.com


Juan Moreno Romo

Support Escalation Engineer Microsoft

- **Based in Madrid, Spain**
 - 9 Years @ Microsoft
 - **Supporting SQL Server On-Premises**
 - SQL Server Core team
 - HA SME
-
- **Speaker @ SQL Saturday, SQL Day, Azure Global Bootcamp, Summit, NetCoreConf, CodeCamp, etc.**
 - Developer, DBA, Data Modeler, Software Architect, IT Manager, Project Manager, SQL Server and SQL Server on Azure VM Support Escalation Engineer.



 /JuanEMorenoRomo

 @DarthSicuel

 Juan Moreno Romo

 jumoreno@microsoft.com

- **Welcome & Expectations (x(15) min)**
- **Introduction to Query Store & Copilot (x(30) min)**
- **Query Store Internals & Diagnostics (x(40) min)**
- **Hands-on Scenarios with Python Simulation (x(90) min)**
- **Support Insights & Best Practices (x(20) min)**
- **Q&A + Wrap Up (x(15) min)**



Welcome & Expectations

- Introduction of the speaker
- Workshop objectives
- Overview of the agenda
- **Interactive round:** Each participant briefly shares:
 - Role & experience with SQL Server / Azure SQL
 - Familiarity with Query Store
 - Familiarity with Copilot or AI tools



Session Goals

- **Identify Common Performance Issues:**
 - Recognize typical scenarios that affect application and database performance.
- **Utilize Diagnostic Tools:**
 - **Query Data Store:** Monitor queries and analyze performance patterns.
 - **SSMS Copilot:** Receive recommendations and solutions to detected issues.
 - Using QDS and SSMS Copilot to prevent the issues to happen
- **Apply Best Coding Practices:**
 - Optimize SQL queries within Python code.
 - Implement strategies to avoid locks and deadlocks.
- **Foster Effective Collaboration:**
 - Understand how developers and DBAs can work together to solve problems.
 - Establish communication channels to prevent future issues.



Session Goals

- **Resolve Problems in Real-Time:**
 - Address common errors and exceptions in Python applications connected to databases.
 - Implement retry techniques and secure transaction handling.
- **Enhance Overall Application Performance:**
 - Apply optimization techniques that benefit both the code and the database.
 - Comprehend the impact of design decisions on performance.
- **Key Message: Empower Yourself to Diagnose and Resolve Issues:**
 - The session will equip you with the skills and knowledge to proactively tackle performance problems.



Connection Details

Server:

dataconwestus2.database.windows.net

User & Password (SQL Login)

Assigned in the post-it that you have.

Database:

PerfTroubleshootingDB

SLO:

Serverless, Gen5, 4 vCores

Introduction to Query Store & Copilot

- **What** is Query Store and **why it matters**
- **Key differences:** On-prem vs Azure SQL
- Introduction to Azure SQL Copilot and SSMS Copilot
- **Demo:** Ask Copilot to analyze query performance from yesterday

Query Store Internals & Diagnostics

- Understanding **query_id, plan_id, runtime stats, wait stats**
- Detecting **regressions, most expensive queries, and anomalies**
- Forcing plans: **when and why**
- Query Store across **multiple databases**
- *Demo:* **Investigate a degraded query using SSMS or Azure Copilot**

Hands-on Scenarios with Python Simulation

- **High CPU** (single and multi-thread)
- **CXPACKET** (query parallel execution)
- **Different Execution Plan** (parameter sniffing)
- **Command Timeout** with Retry Logic
 - **ODBC Connection Reuse after Error**
- **Simulation High Data Read and Write.**
- **High Concurrency**
- **Additional Topic - Coding Stuff.**

Python Code



Session Context:

Introduction to a Python application that simulates common database bad performance scenarios.

Focus on identifying and resolving these issues collaboratively.




Tools We'll Use:

Python: The programming language used to develop the application and simulate scenarios.

Query Data Store (QDS): A tool for monitoring and analyzing query performance in SQL Server.

SQL Copilot: An AI-powered assistant that helps diagnose and solve database issues.



Importance of the Topic:

Performance issues can significantly impact applications and user experience.

Collaboration between developers and DBAs is crucial for efficiently identifying and resolving these problems.



Expectations for Attendees:

No prior experience with the mentioned tools is required.

The session will be practical and focused on real-world examples.

Participation and questions are encouraged throughout the presentation.

Python Connection Details

Authentication:

ODBC Driver 18

username/password (via credentials.txt)

Connection Handling:

Defined in ConnectToTheDB() in Python code.

Supports retry, timeout, and logging

Tables involved

Schema	Table	Description & Usage
Dbo	Inventory	Used for high concurrency simulation with locking
Dbo	MS_TableA_MS / MS_TableB_MS	Used in deadlock simulation (simulate_deadlock)
Dbo	Notes	Accessed via GiveNotes for plan regression testing
Dbo	Products2	Queried in chatty app simulation with optional caching
MSxyzTest	_x_y_z_MS_HighAsyncNetworkIO	Used to simulate network latency issues
MSxyzTest	_x_y_z_MS_HighBulkInsert	Prepared for bulk insert scenarios
MSxyzTest	_x_y_z_MS_HighCPU	Main target for High CPU test (looping query with TextToSearch)
MSxyzTest	_x_y_z_MS_HighCXPacket	Used for CXPACKET scenarios with complex sort/filter logic
MSxyzTest	_x_y_z_MS_HighDATAIO / HighDATAIOBlocks	Referenced in inefficient data access patterns
MSxyzTest	_x_y_z_MS_HighLocks	Included for potential future use in lock contention scenarios
MSxyzTest	_x_y_z_MS_HighLogIO	Placeholder for log IO stress scenarios
MSxyzTest	_x_y_z_MS_HighTempDB	Used in tempdb contention simulations


Python Code Modules Overview (Database Workload)

- **RunHighCPU()**
 - Executes many CPU-intensive queries using VARCHAR or WCHAR
- **RunDifferentExecutionPlan()**
 - Calls stored procedure with random inputs to vary plan
- **RunCommandTimeout()**
 - Demonstrates timeout + retry logic pattern
- **RunHighNetworkIO**
 - High Network Latency and DataIO Write (PowerShell)
- **RunConcurrency()**
 - A lot of blocking issues.

Hands-On-Lab I

High CPU and CXPacket

High CPU Usage

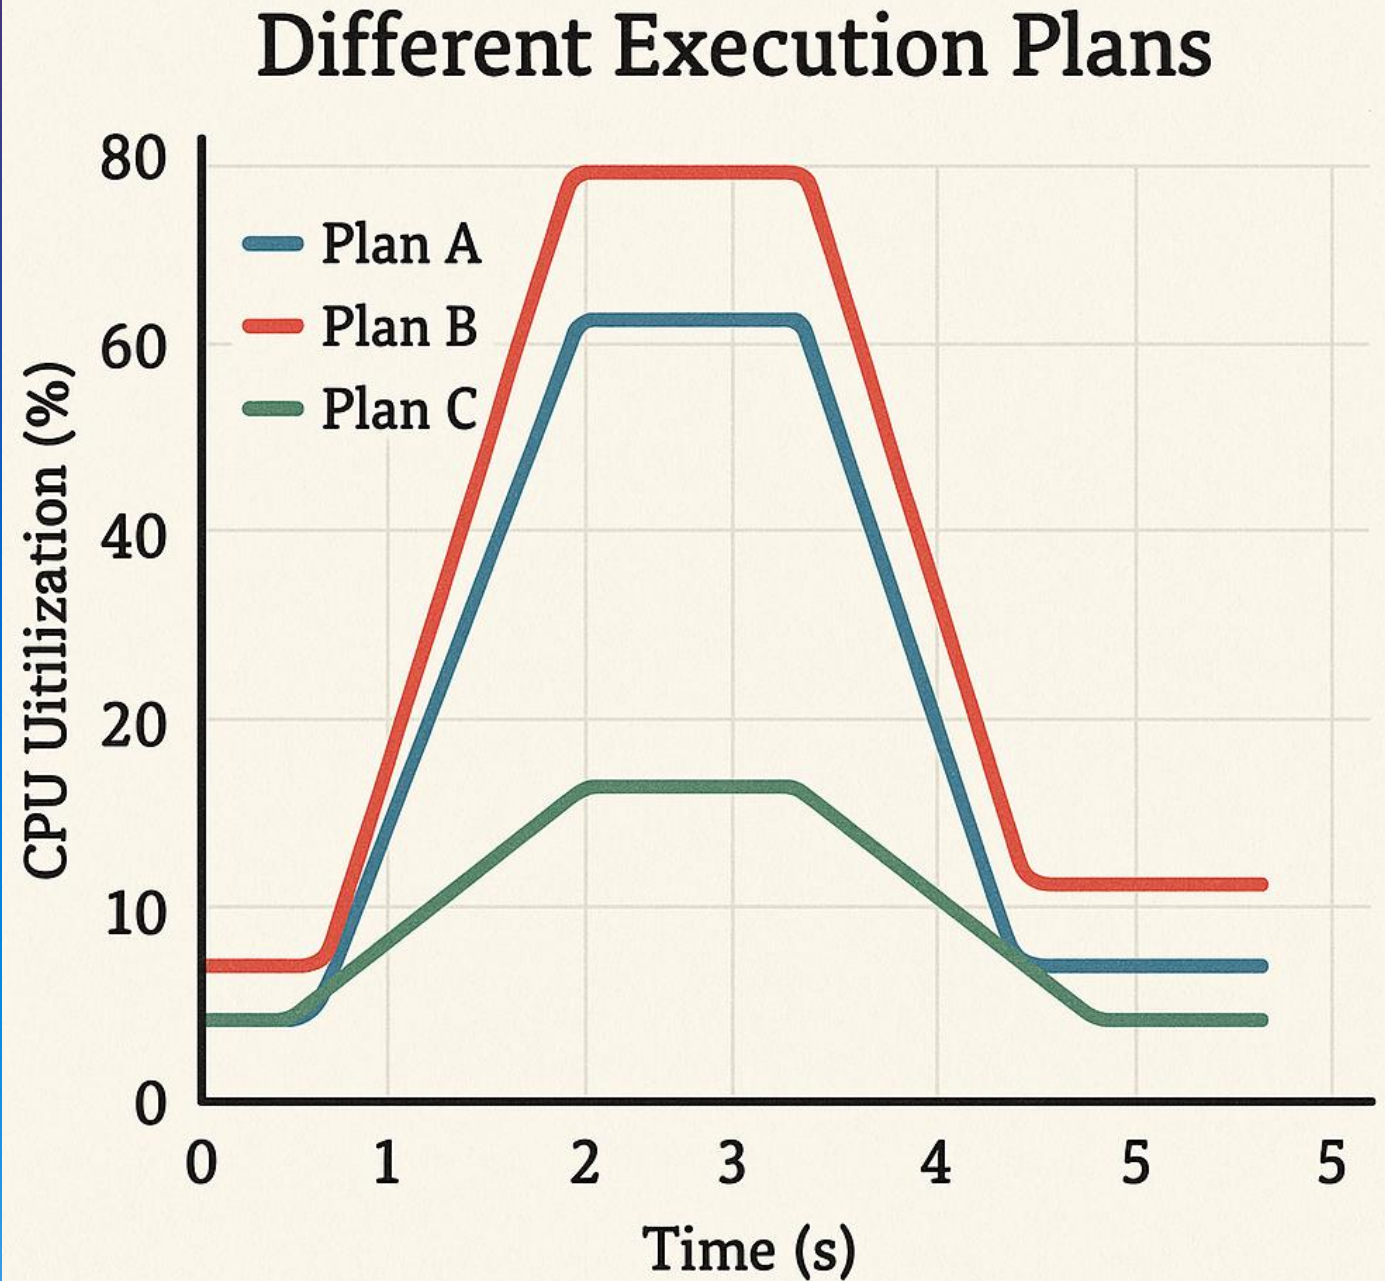


92%
CPU usage

[View details](#)

Hands-On-Lab II

Different Execution Plans




Hands-On-Lab III

Command Timeout vs
Connection Timeout



Hands-On-Lab IV

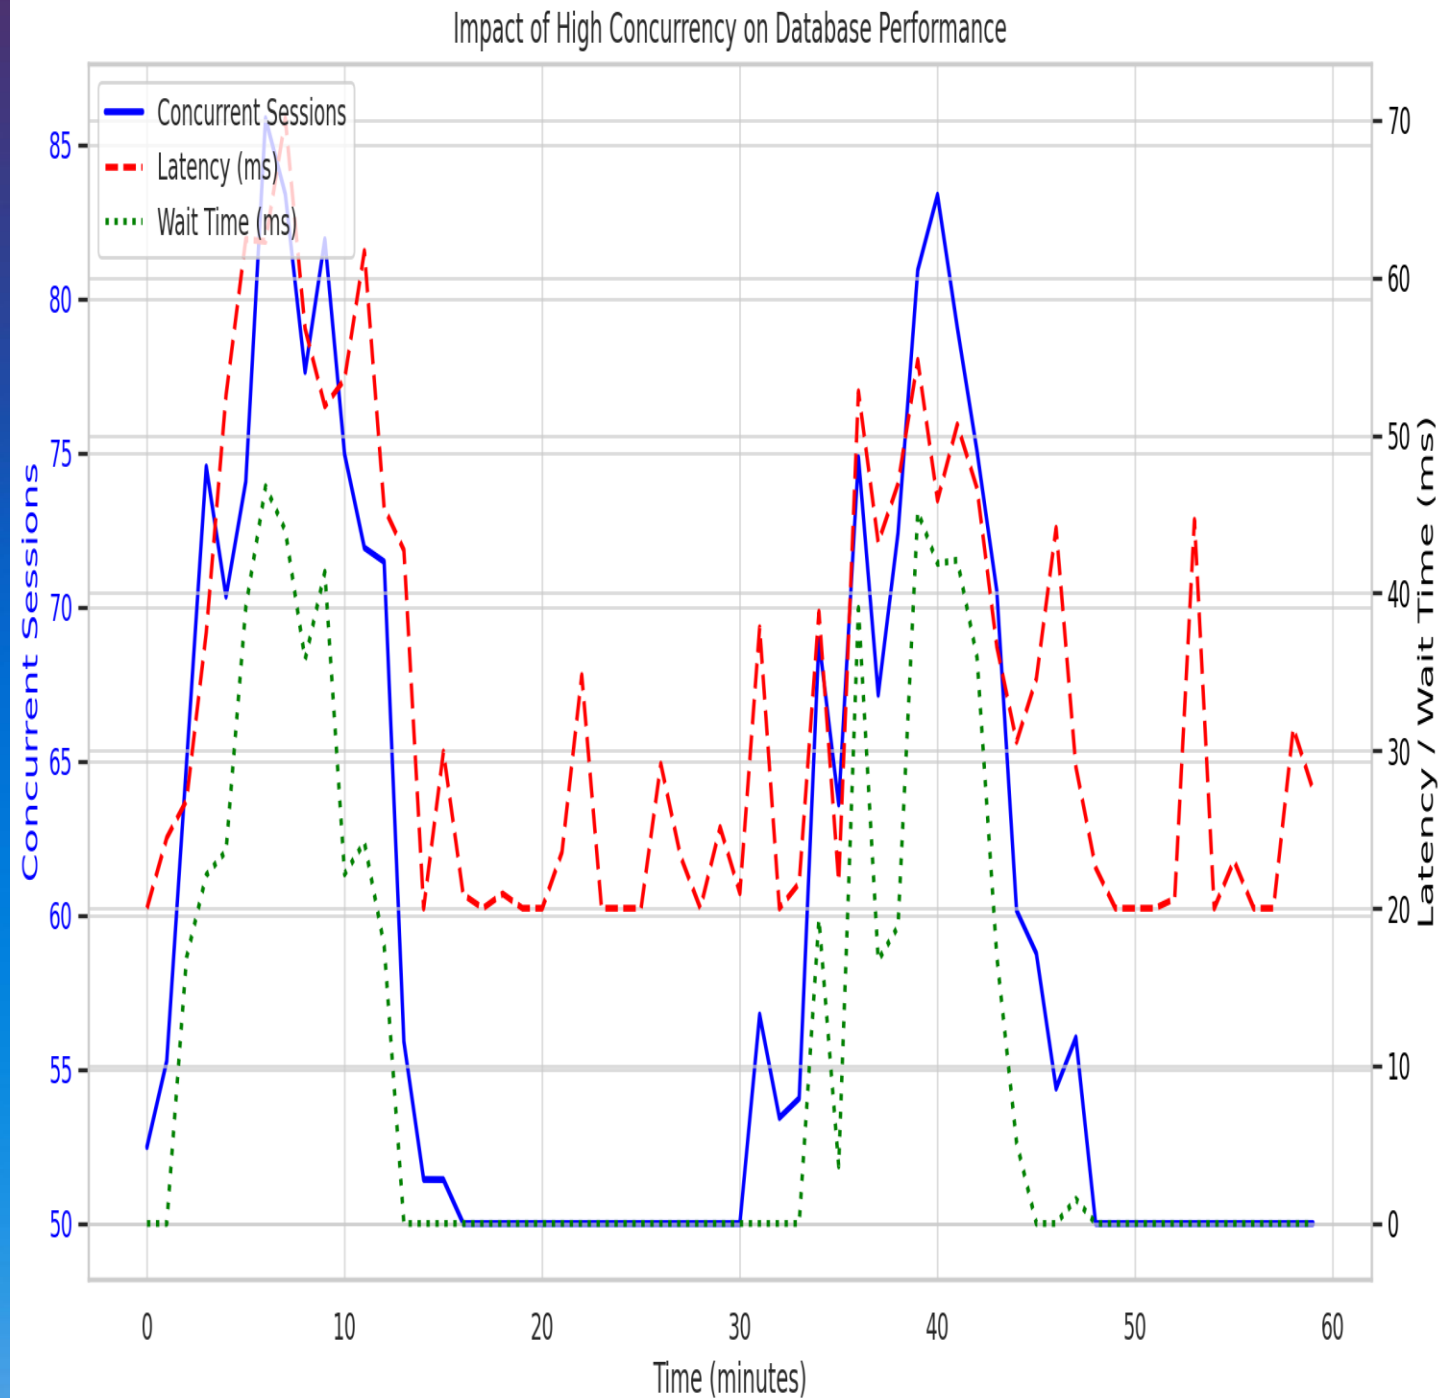
PowerShell Script bulkInsert and
DataIO



**HIGH
READ/WRITE I/O**

Hands-On-Lab V

Concurrency



Hands-On-Lab VI

Finding Performance Issue
(Portal, QDS and Copilot).



Hands-On-Lab VII

QDS Summary Across Elastic
database Pool





Support Insights & Best Practices

- **Real-world support** case learnings
- Query Store for **proactive health monitoring**
- When to use **Copilot vs traditional tools**

Python Code Modules Overview

(Coding Best Practices)

- **ConnectToTheDB()**
 - Central function for all DB connections, manages retries and logs
- **run_connection_benchmark()**
 - Stress test for connection time and limits
- **RunSimpleInefficiencyWithTiming()**
 - Efficient vs inefficient data processing example
- **create_guiNumberOfExecutions()**
 - Stress test for execution number of a limited of threads.

Q&A + Wrap Up

- **Final** questions
- **Sharing resources and GitHub repo**
- What will you apply after this session?



Thank you.

Jose Manuel Jurado Diaz

Escalation Engineer

Microsoft

Juan Moreno Romo

Support Escalation Engineer

Microsoft

Queries right now? *Shoot them!*

Queries in future? *Ping them!*

- How can I obtain the query with ID 1?
- Which queries are consuming the most CPU in the database in the last hour?
- Which queryID Text are consuming the most CPU in the database in the last hour, show me the query text?
- How can I optimize the query with ID 4?
- How can I identify and add missing indexes for my database?
- Are there any missing index suggestions for the high CPU-consuming queries?
- Which queries are being canceled due to timeouts?
- Which queries are exhibiting common performance antipatterns?
- What is the average and maximum duration of the queries canceled by timeouts?
- Which queries are using inefficient joins or subqueries?
- Can you provide a list of missing indexes in my database?
- What are the most common wait times in my database?
- Which queries are using forced execution plans?
- Is there any command timeout in the database?
- Is there any query with different execution plans?
- What are the reasons for multiple execution plans being generated for the same query?
- Is parameter sniffing causing performance issues with queries that have multiple execution plans?
- give the TSQL command statement to identify the query_id where the sql statement is "SELECT count(Name),name FROM Notes where ID<@n group by Name"
- give the TSQL command statement to identify the query_id where the sql statement is like "SELECT count(Name),name FROM Notes where ID<@n group by Name"
- How many times has the query with id 4 been executed?
- Show me the execution history of query ID 4.
- Can you show the execution plan for query ID 4?
- Are there any queries using unnecessary SELECT * operations?
- Can you provide a detailed list of wait times and their deltas?
- Can you provide the text of the queries that are consuming the most CPU?